

Methods for solving Linear Least Squares problems

Anibal Sosa



IPM for Linear Programming,
September 2009



Outline

- 1 The Least Square Problem (LSQ)
 - Linear Least Square Problems
- 2 Methods for solving Linear LSQ
 - Normal Equations
 - QR Factorization
 - Singular Value Decomposition (SVD)
- 3 Comments on the three methods
- 4 Regularization techniques
 - Tikhonov regularization and Damped SVD
 - Tikhonov regularization order one and two

Outline

- 1 The Least Square Problem (LSQ)
 - Linear Least Square Problems
- 2 Methods for solving Linear LSQ
 - Normal Equations
 - QR Factorization
 - Singular Value Decomposition (SVD)
- 3 Comments on the three methods
- 4 Regularization techniques
 - Tikhonov regularization and Damped SVD
 - Tikhonov regularization order one and two

Outline

- 1 The Least Square Problem (LSQ)
 - Linear Least Square Problems
- 2 Methods for solving Linear LSQ
 - Normal Equations
 - QR Factorization
 - Singular Value Decomposition (SVD)
- 3 Comments on the three methods
- 4 Regularization techniques
 - Tikhonov regularization and Damped SVD
 - Tikhonov regularization order one and two

Outline

- 1 The Least Square Problem (LSQ)
 - Linear Least Square Problems
- 2 Methods for solving Linear LSQ
 - Normal Equations
 - QR Factorization
 - Singular Value Decomposition (SVD)
- 3 Comments on the three methods
- 4 Regularization techniques
 - Tikhonov regularization and Damped SVD
 - Tikhonov regularization order one and two

The Least Square Problem (LSQ)

- The objective function has the following special form

$$f(x) = \frac{1}{2} \sum_{j=1}^m r_j^2(x), \text{ where } r_j : \mathbb{R}^n \rightarrow \mathbb{R} \text{ are the } \textit{residuals}, \text{ i. e.,}$$

$$\min_{x \in \mathbb{R}^n} f(x) = \min_{x \in \mathbb{R}^n} \frac{1}{2} r^T(x) r(x) = \min_{x \in \mathbb{R}^n} \frac{1}{2} \|r(x)\|_2^2$$

$$r : \mathbb{R}^n \rightarrow \mathbb{R}^m \text{ is called the residual vector, i.e., } r = \begin{bmatrix} r_1(x) \\ r_2(x) \\ \vdots \\ r_m(x) \end{bmatrix}$$

- Least square problems arise in many areas of applications
- Largest source of unconstrained optimization problems

The Least Square Problem (LSQ)

- The objective function has the following special form

$$f(x) = \frac{1}{2} \sum_{j=1}^m r_j^2(x), \text{ where } r_j : \mathbb{R}^n \rightarrow \mathbb{R} \text{ are the } \textit{residuals}, \text{ i. e.,}$$

$$\min_{x \in \mathbb{R}^n} f(x) = \min_{x \in \mathbb{R}^n} \frac{1}{2} r^T(x) r(x) = \min_{x \in \mathbb{R}^n} \frac{1}{2} \|r(x)\|_2^2$$

$$r : \mathbb{R}^n \rightarrow \mathbb{R}^m \text{ is called the residual vector, i.e., } r = \begin{bmatrix} r_1(x) \\ r_2(x) \\ \vdots \\ r_m(x) \end{bmatrix}$$

- Least square problems arise in many areas of applications
- Largest source of unconstrained optimization problems

The Least Square Problem (LSQ)

- The objective function has the following special form

$$f(x) = \frac{1}{2} \sum_{j=1}^m r_j^2(x), \text{ where } r_j : \mathbb{R}^n \rightarrow \mathbb{R} \text{ are the } \textit{residuals} , \text{ i. e.,}$$

$$\min_{x \in \mathbb{R}^n} f(x) = \min_{x \in \mathbb{R}^n} \frac{1}{2} r^T(x) r(x) = \min_{x \in \mathbb{R}^n} \frac{1}{2} \|r(x)\|_2^2$$

$$r : \mathbb{R}^n \rightarrow \mathbb{R}^m \text{ is called the residual vector, i.e., } r = \begin{bmatrix} r_1(x) \\ r_2(x) \\ \vdots \\ r_m(x) \end{bmatrix}$$

- Least square problems arise in many areas of applications
- Largest source of unconstrained optimization problems

Linear Least Square Problems

- Let $\phi(x; \rho)$ be a model function that predict experimental values, for some fix parameters ρ .
- Usually we want to minimize the differences between the observed values $y \in \mathbb{R}^m$ (data) and the predicted values $\phi(x; \rho) \in \mathbb{R}^m$.
- We can use LSQ setting $r(x) = \phi(x; \rho) - y$

$$\min_{x \in \mathbb{R}^n} \frac{1}{2} \|\phi(x; \rho) - y\|_2^2 \quad (1)$$

- If ϕ in (2) is nonlinear then we have a nonlinear LSQ problem
- In our case $\phi(x) = Ax$, thus we say this is a linear LSQ problem

Linear Least Square Problems

- Let $\phi(x; \rho)$ be a model function that predict experimental values, for some fix parameters ρ .
- Usually we want to minimize the differences between the observed values $y \in \mathbb{R}^m$ (data) and the predicted values $\phi(x; \rho) \in \mathbb{R}^m$.
- We can use LSQ setting $r(x) = \phi(x; \rho) - y$

$$\min_{x \in \mathbb{R}^n} \frac{1}{2} \|\phi(x; \rho) - y\|_2^2 \quad (1)$$

- If ϕ in (2) is nonlinear then we have a nonlinear LSQ problem
- In our case $\phi(x) = Ax$, thus we say this is a linear LSQ problem

Linear Least Square Problems

- Let $\phi(x; \rho)$ be a model function that predict experimental values, for some fix parameters ρ .
- Usually we want to minimize the differences between the observed values $y \in \mathbb{R}^m$ (data) and the predicted values $\phi(x; \rho) \in \mathbb{R}^m$.
- We can use LSQ setting $r(x) = \phi(x; \rho) - y$

$$\min_{x \in \mathbb{R}^n} \frac{1}{2} \|\phi(x; \rho) - y\|_2^2 \quad (1)$$

- If ϕ in (2) is nonlinear then we have a nonlinear LSQ problem
- In our case $\phi(x) = Ax$, thus we say this is a linear LSQ problem

Linear Least Square Problems

- Let $\phi(x; \rho)$ be a model function that predict experimental values, for some fix parameters ρ .
- Usually we want to minimize the differences between the observed values $y \in \mathbb{R}^m$ (data) and the predicted values $\phi(x; \rho) \in \mathbb{R}^m$.
- We can use LSQ setting $r(x) = \phi(x; \rho) - y$

$$\min_{x \in \mathbb{R}^n} \frac{1}{2} \|\phi(x; \rho) - y\|_2^2 \quad (1)$$

- If ϕ in (2) is nonlinear then we have a nonlinear LSQ problem
- In our case $\phi(x) = Ax$, thus we say this is a linear LSQ problem

Preliminaries for solving the LSQ problem

Observe that

$$f(x) = \frac{1}{2} \|Ax - y\|_2^2 = \frac{1}{2} (Ax - y)^T (Ax - y) = \frac{1}{2} x^T A^T A x - x^T A^T y + \frac{1}{2} y^T y$$

is easy to prove that

$$\nabla f(x) = A^T (Ax - y) \quad \nabla^2 f(x) = A^T A$$

Since f is a convex function is well known that any x^* such that $\nabla f(x^*) = 0$ is a global minimizer of f , therefore x^* satisfy the normal equations

$$A^T A x = A^T y$$

Next we discuss three major algorithms for solving Linear LSQ problems, assuming: i) $m \geq n$ and ii) A is full rank

Preliminaries for solving the LSQ problem

Observe that

$$f(x) = \frac{1}{2} \|Ax - y\|_2^2 = \frac{1}{2} (Ax - y)^T (Ax - y) = \frac{1}{2} x^T A^T A x - x^T A^T y + \frac{1}{2} y^T y$$

is easy to prove that

$$\nabla f(x) = A^T (Ax - y) \quad \nabla^2 f(x) = A^T A$$

Since f is a convex function is well known that any x^* such that $\nabla f(x^*) = 0$ is a global minimizer of f , therefore x^* satisfy the normal equations

$$A^T A x = A^T y$$

Next we discuss three major algorithms for solving Linear LSQ problems, assuming: i) $m \geq n$ and ii) A is full rank

Preliminaries for solving the LSQ problem

Observe that

$$f(x) = \frac{1}{2} \|Ax - y\|_2^2 = \frac{1}{2} (Ax - y)^T (Ax - y) = \frac{1}{2} x^T A^T A x - x^T A^T y + \frac{1}{2} y^T y$$

is easy to prove that

$$\nabla f(x) = A^T (Ax - y) \quad \nabla^2 f(x) = A^T A$$

Since f is a convex function is well known that any x^* such that $\nabla f(x^*) = 0$ is a global minimizer of f , therefore x^* satisfy the normal equations

$$A^T A x = A^T y$$

Next we discuss three major algorithms for solving Linear LSQ problems, assuming: i) $m \geq n$ and ii) A is full rank

Normal Equations

Step1 Compute $A^T A$ and $A^T y$

Step2 Compute Cholesky factorization of $A^T A > 0$

$$A^T A = R^T R, \quad R \text{ is an upper triangular matrix} (R_{ii} > 0)$$

Step3 Perform two triangular substitutions

$$R^T z = R^T y \implies R x^* = z$$

Disadvantages:

- Relative error of $x^* \approx \kappa(A)^{21}$
- Sensitive to ill-conditioned matrices

$$\kappa(A) = \|A\| \|A^{-1}\| \approx \frac{\sigma_1}{\sigma_n} = \kappa_2(A)$$

Normal Equations

Step1 Compute $A^T A$ and $A^T y$

Step2 Compute Cholesky factorization of $A^T A > 0$

$$A^T A = R^T R, \quad R \text{ is an upper triangular matrix} (R_{ii} > 0)$$

Step3 Perform two triangular substitutions

$$R^T z = R^T y \implies R x^* = z$$

Disadvantages:

- Relative error of $x^* \approx \kappa(A)^{21}$
- Sensitive to ill-conditioned matrices

$$\kappa(A) = \|A\| \|A^{-1}\| \approx \frac{\sigma_1}{\sigma_n} = \kappa_2(A)$$

Normal Equations

Step1 Compute $A^T A$ and $A^T y$

Step2 Compute Cholesky factorization of $A^T A > 0$

$$A^T A = R^T R, \quad R \text{ is an upper triangular matrix } (R_{ii} > 0)$$

Step3 Perform two triangular substitutions

$$R^T z = R^T y \implies R x^* = z$$

Disadvantages:

- Relative error of $x^* \approx \kappa(A)^{21}$
- Sensitive to ill-conditioned matrices

$$^1 \kappa(A) = \|A\| \|A^{-1}\| \approx \frac{\sigma_1}{\sigma_n} = \kappa_2(A)$$

QR Factorization

Notice that $\|\cdot\|$ is invariant under orthogonal transformations

$$\|Ax - y\|_2^2 = \|Q^T(Ax - y)\|_2^2$$

where $Q_{m \times m}$ is orthogonal

- The QR factorization is done as follows

$$A\Pi = Q \begin{bmatrix} R \\ 0 \end{bmatrix} = [Q_1 \quad Q_2] \begin{bmatrix} R \\ 0 \end{bmatrix} = Q_1 R \quad (2)$$

where $\Pi_{n \times n}$ is a permutation matrix, Q_1 is the first n columns of Q and $R_{n \times n}$ is upper triangular with $R_{ii} > 0$

- Using 2 we have

$$\|Ax - y\|_2^2 = \left\| \begin{bmatrix} Q_1^T \\ Q_2^T \end{bmatrix} (A\Pi\Pi^T x - y) \right\|_2^2$$

QR Factorization

Notice that $\|\cdot\|$ is invariant under orthogonal transformations

$$\|Ax - y\|_2^2 = \|Q^T(Ax - y)\|_2^2$$

where $Q_{m \times m}$ is orthogonal

- The QR factorization is done as follows

$$A\Pi = Q \begin{bmatrix} R \\ 0 \end{bmatrix} = [Q_1 \quad Q_2] \begin{bmatrix} R \\ 0 \end{bmatrix} = Q_1 R \quad (2)$$

where $\Pi_{n \times n}$ is a permutation matrix, Q_1 is the first n columns of Q and $R_{n \times n}$ is upper triangular with $R_{ii} > 0$

- Using 2 we have

$$\|Ax - y\|_2^2 = \left\| \begin{bmatrix} Q_1^T \\ Q_2^T \end{bmatrix} (A\Pi\Pi^T x - y) \right\|_2^2$$

QR Factorization

Notice that $\|\cdot\|$ is invariant under orthogonal transformations

$$\|Ax - y\|_2^2 = \|Q^T(Ax - y)\|_2^2$$

where $Q_{m \times m}$ is orthogonal

- The QR factorization is done as follows

$$A\Pi = Q \begin{bmatrix} R \\ 0 \end{bmatrix} = [Q_1 \quad Q_2] \begin{bmatrix} R \\ 0 \end{bmatrix} = Q_1 R \quad (2)$$

where $\Pi_{n \times n}$ is a permutation matrix, Q_1 is the first n columns of Q and $R_{n \times n}$ is upper triangular with $R_{ii} > 0$

- Using 2 we have

$$\|Ax - y\|_2^2 = \left\| \begin{bmatrix} Q_1^T \\ Q_2^T \end{bmatrix} (A\Pi\Pi^T x - y) \right\|_2^2$$



QR Factorization(2)

$$\left| \begin{bmatrix} Q_1^T \\ Q_2^T \end{bmatrix} \left(\underbrace{[Q_1 \ Q_2]}_{A\Pi} \begin{bmatrix} R \\ 0 \end{bmatrix} \Pi^T x - y \right) \right|_2^2 = \left| \begin{bmatrix} R \\ 0 \end{bmatrix} \Pi^T x - \begin{bmatrix} Q_1^T \\ Q_2^T \end{bmatrix} y \right|_2^2$$

$$= \|R\Pi^T x - Q_1^T y\|_2^2 + \|Q_2^T y\|_2^2$$

Notice that from the last equation:

- The last term does not depend on x
- The minimum value is reached when $R\Pi^T x - Q_1^T y = 0$, therefore

$$x^* = \Pi R^{-1} Q_1^T y$$

QR Factorization(2)

$$\left| \begin{bmatrix} Q_1^T \\ Q_2^T \end{bmatrix} \left(\underbrace{[Q_1 \ Q_2]}_{A\Pi} \begin{bmatrix} R \\ 0 \end{bmatrix} \Pi^T x - y \right) \right|_2^2 = \left| \begin{bmatrix} R \\ 0 \end{bmatrix} \Pi^T x - \begin{bmatrix} Q_1^T \\ Q_2^T \end{bmatrix} y \right|_2^2$$

$$= \|R\Pi^T x - Q_1^T y\|_2^2 + \|Q_2^T y\|_2^2$$

Notice that from the last equation:

- The last term does not depend on x
- The minimum value is reached when $R\Pi^T x - Q_1^T y = 0$, therefore

$$x^* = \Pi R^{-1} Q_1^T y$$

QR Factorization Algorithm

Step1 Compute QR factorization of A

Step2 Extract Q_1 , identify Π and R

Step3 Perform one triangular substitution and one permutation

$$Rz = Q_1^T y \implies x^* = \Pi z$$

Advantage:

- Relative error of $x^* \approx \kappa(A)$

Disadvantage:

- Sometimes is necessary more information about data sensitivity

QR Factorization Algorithm

Step1 Compute QR factorization of A

Step2 Extract Q_1 , identify Π and R

Step3 Perform one triangular substitution and one permutation

$$Rz = Q_1^T y \implies x^* = \Pi z$$

Advantage:

- Relative error of $x^* \approx \kappa(A)$

Disadvantage:

- Sometimes is necessary more information about data sensitivity

QR Factorization Algorithm

Step1 Compute QR factorization of A

Step2 Extract Q_1 , identify Π and R

Step3 Perform one triangular substitution and one permutation

$$Rz = Q_1^T y \implies x^* = \Pi z$$

Advantage:

- Relative error of $x^* \approx \kappa(A)$

Disadvantage:

- Sometimes is necessary more information about data sensitivity

Singular Value Decomposition (SVD)

Theorem

If $A_{m \times n}$ is real then there exist orthogonal matrices

$$U = [u_1 \dots u_m] \in \mathbb{R}^{m \times m} \text{ and } V = [v_1 \dots v_n] \in \mathbb{R}^{n \times n}$$

such that $A = U\Sigma V^T$, where

$$\Sigma = \text{diag}(\sigma_1, \dots, \sigma_p) \in \mathbb{R}^{m \times n}, \quad p = \min\{m, n\} \text{ and } \sigma_1 \geq \sigma_2 \dots \geq \sigma_p \geq 0$$

- In our case $\sigma_1 \geq \sigma_2 \dots \geq \sigma_n > 0$ since A is full rank and $m \gg n$ thus

$$A = U \begin{bmatrix} \Sigma_1 \\ 0 \end{bmatrix} V^T = [U_1 \quad U_2] \begin{bmatrix} \Sigma_1 \\ 0 \end{bmatrix} V^T = U_1 \Sigma_1 V^T \quad (3)$$

where U_1 has the first n columns of U and $\Sigma_1 = \text{diag}(\sigma_1, \dots, \sigma_n)$.

Singular Value Decomposition (SVD)

Theorem

If $A_{m \times n}$ is real then there exist orthogonal matrices

$$U = [u_1 \dots u_m] \in \mathbb{R}^{m \times m} \text{ and } V = [v_1 \dots v_n] \in \mathbb{R}^{n \times n}$$

such that $A = U \Sigma V^T$, where

$$\Sigma = \text{diag}(\sigma_1, \dots, \sigma_p) \in \mathbb{R}^{m \times n}, \quad p = \min\{m, n\} \text{ and } \sigma_1 \geq \sigma_2 \dots \geq \sigma_p \geq 0$$

- In our case $\sigma_1 \geq \sigma_2 \dots \geq \sigma_n > 0$ since A is full rank and $m \gg n$ thus

$$A = U \begin{bmatrix} \Sigma_1 \\ 0 \end{bmatrix} V^T = [U_1 \quad U_2] \begin{bmatrix} \Sigma_1 \\ 0 \end{bmatrix} V^T = U_1 \Sigma_1 V^T \quad (3)$$

where U_1 has the first n columns of U and $\Sigma_1 = \text{diag}(\sigma_1, \dots, \sigma_n)$.

The thin SVD

- Using (3) and similar ideas from QR

$$\begin{aligned}\|Ax - y\|_2^2 &= \left\| \begin{bmatrix} \Sigma_1 \\ 0 \end{bmatrix} (V^T x) - \begin{bmatrix} U_1^T \\ U_2^T \end{bmatrix} y \right\|_2^2 \\ &= \|\Sigma_1 (V^T x) - U_1^T y\|_2^2 + \|U_2^T y\|_2^2\end{aligned}$$

Again from the last equation:

- The last term does not depend on x
- The minimum value is reached when $\Sigma (V^T x) - U_1^T y = 0$, therefore

$$x^* = V \Sigma^{-1} U_1^T y$$

or equivalently

$$x^* = \sum_{i=1}^n \left(\frac{u_i^T y}{\sigma_i} \right) v_i$$



The thin SVD

- Using (3) and similar ideas from QR

$$\begin{aligned}\|Ax - y\|_2^2 &= \left\| \begin{bmatrix} \Sigma_1 \\ 0 \end{bmatrix} (V^T x) - \begin{bmatrix} U_1^T \\ U_2^T \end{bmatrix} y \right\|_2^2 \\ &= \|\Sigma_1 (V^T x) - U_1^T y\|_2^2 + \|U_2^T y\|_2^2\end{aligned}$$


Again from the last equation:

- The last term does not depend on x
- The minimum value is reached when $\Sigma (V^T x) - U_1^T y = 0$, therefore

$$x^* = V \Sigma^{-1} U_1^T y$$

or equivalently

$$x^* = \sum_{i=1}^n \left(\frac{u_i^T y}{\sigma_i} \right) v_i$$



SVD

- Equation (4) gives useful information about x^* sensitivity
 - Small changes in A or y can induce large changes in x^* if σ_i is small
 - A is rank deficient when $\frac{\sigma_n}{\sigma_1} \ll 1$. (σ_n is the distance from A to the set of singular matrices)
- x^* calculated as in (4) has the smallest 2-norm of all minimizers

Advantage:

- Most robust and reliable

Disadvantage:

- Most expensive

SVD

- Equation (4) gives useful information about x^* sensitivity
 - Small changes in A or y can induce large changes in x^* if σ_i is small
 - A is rank deficient when $\frac{\sigma_n}{\sigma_1} \ll 1$. (σ_n is the distance from A to the set of singular matrices)
- x^* calculated as in (4) has the smallest 2-norm of all minimizers

Advantage:

- Most robust and reliable

Disadvantage:

- Most expensive

Normal Eq. vs QR vs SVD

- The Cholesky-based algorithm is practical if $m \gg n$ (is easier store $A^T A$), even if A is sparse
- The QR algorithm avoid squaring $\kappa(A)$
- When A is rank-deficient, some $\sigma_i \approx 0$ thus any vector

$$x^* = \sum_{\sigma_i \neq 0} \left(\frac{u_i^T y}{\sigma_i} \right) v_i + \sum_{\sigma_i = 0} \tau v_i$$

is also a minimizer of $\|Ax - y\|$, for τ such that $\sigma_i \geq \tau$. Thus setting $\tau_i = 0$ we get the minimum norm solution²

Remark: For very large problems is recommended to use iterative methods as *Conjugate Gradient*

²This is a type of filter by doing truncation

Normal Eq. vs QR vs SVD

- The Cholesky-based algorithm is practical if $m \gg n$ (is easier store $A^T A$), even if A is sparse
- The QR algorithm avoid squaring $\kappa(A)$
- When A is rank-deficient, some $\sigma_i \approx 0$ thus any vector

$$x^* = \sum_{\sigma_i \neq 0} \left(\frac{u_i^T y}{\sigma_i} \right) v_i + \sum_{\sigma_i = 0} \tau v_i$$

is also a minimizer of $\|Ax - y\|$, for τ such that $\sigma_i \geq \tau$. Thus setting $\tau_i = 0$ we get the minimum norm solution²

Remark: For very large problems is recommended to use iterative methods as *Conjugate Gradient*

²This is a type of filter by doing truncation

Normal Eq. vs QR vs SVD

- The Cholesky-based algorithm is practical if $m \gg n$ (is easier store $A^T A$), even if A is sparse
- The QR algorithm avoid squaring $\kappa(A)$
- When A is rank-deficient, some $\sigma_i \approx 0$ thus any vector

$$x^* = \sum_{\sigma_i \neq 0} \left(\frac{u_i^T y}{\sigma_i} \right) v_i + \sum_{\sigma_i = 0} \tau v_i$$

is also a minimizer of $\|Ax - y\|$, for τ such that $\sigma_i \geq \tau$,. Thus setting $\tau_i = 0$ we get the minimum norm solution²

Remark: For very large problems is recommended to use iterative methods as *Conjugate Gradient*

²This is a type of filter by doing truncation

Normal Eq. vs QR vs SVD

- The Cholesky-based algorithm is practical if $m \gg n$ (is easier store $A^T A$), even if A is sparse
- The QR algorithm avoid squaring $\kappa(A)$
- When A is rank-deficient, some $\sigma_i \approx 0$ thus any vector

$$x^* = \sum_{\sigma_i \neq 0} \left(\frac{u_i^T y}{\sigma_i} \right) v_i + \sum_{\sigma_i = 0} \tau v_i$$

is also a minimizer of $\|Ax - y\|$, for τ such that $\sigma_i \geq \tau$. Thus setting $\tau_i = 0$ we get the minimum norm solution²

Remark: For very large problems is recommended to use iterative methods as *Conjugate Gradient*

²This is a type of filter by doing truncation

Tikhonov regularization^a

^aRidge regression

- Most commonly used method for ill-posed problems
- The ill-conditioned problem 1 is posed as

$$\min \frac{1}{2} \|Ax - y\|_2^2 + \frac{1}{2} \alpha^2 \|x\|_2^2 \quad (5)$$

for some suitable *regularization parameter* $\alpha > 0$

- This improves the problem condition, even is A is rank-deficient, shifting the small singular values

$$(A^T A + \alpha I_n) x = \underbrace{A^T A x}_{\lambda x} + \alpha x = (\lambda + \alpha) x$$

for any eigenvalue λ and eigenvector x of $A^T A$

Tikhonov regularization^a

^aRidge regression

- Most commonly used method for ill-posed problems
- The ill-conditioned problem 1 is posed as

$$\min \frac{1}{2} \|Ax - y\|_2^2 + \frac{1}{2} \alpha^2 \|x\|_2^2 \quad (5)$$

for some suitable *regularization parameter* $\alpha > 0$

- This improves the problem condition, even if A is rank-deficient, shifting the small singular values

$$(A^T A + \alpha I_n) x = \underbrace{A^T A x}_{\lambda x} + \alpha x = (\lambda + \alpha) x$$

for any eigenvalue λ and eigenvector x of $A^T A$

Tikhonov regularization^a

^aRidge regression

- Most commonly used method for ill-posed problems
- The ill-conditioned problem 1 is posed as

$$\min \frac{1}{2} \|Ax - y\|_2^2 + \frac{1}{2} \alpha^2 \|x\|_2^2 \quad (5)$$

for some suitable *regularization parameter* $\alpha > 0$

- This improves the problem condition, even if A is rank-deficient, shifting the small singular values

$$(A^T A + \alpha I_n) x = \underbrace{A^T A x}_{\lambda x} + \alpha x = (\lambda + \alpha) x$$

for any eigenvalue λ and eigenvector x of $A^T A$

Tikhonov regularization and Damped SVD

- A little algebra shows that the minimum solution of (5) is given by the nonsingular system

$$(A^T A + \alpha^2 I_n) x = A^T y$$

and from (4) we can show that

$$x^* = \sum_{i=1}^n f_i \left(\frac{u_i^T y}{\sigma_i} \right) v_i$$

where $f_i = \frac{\sigma_i^2}{\sigma_i^2 + \alpha^2}$ are known as *filter factors*³

- The impact of an small α in the filter factors is:
 - None for large σ_i ($\alpha \ll \sigma_i$), i.e. $\frac{\sigma_i^2}{\sigma_i^2 + \alpha^2} \approx 1$
 - Reduce the magnification of $\frac{1}{\sigma_i}$ since $\frac{\sigma_i^2}{\sigma_i^2 + \alpha^2} \approx \frac{\sigma_i^2}{\alpha^2} \ll 1$
- A “good” choice of α may provide enough numerical stability to expect a good approximate solution

³ In singular value decomposition, the singular values σ_i are the square roots of the eigenvalues of $A^T A$.

Tikhonov regularization and Damped SVD

- A little algebra shows that the minimum solution of (5) is given by the nonsingular system

$$(A^T A + \alpha^2 I_n) x = A^T y$$

and from (4) we can show that

$$x^* = \sum_{i=1}^n f_i \left(\frac{u_i^T y}{\sigma_i} \right) v_i$$

where $f_i = \frac{\sigma_i^2}{\sigma_i^2 + \alpha^2}$ are known as *filter factors*³

- The impact of an small α in the filter factors is:
 - None for large σ_i ($\alpha \ll \sigma_i$), i.e. $\frac{\sigma_i^2}{\sigma_i^2 + \alpha^2} \approx 1$
 - Reduce the magnification of $\frac{1}{\sigma_i}$ since $\frac{\sigma_i^2}{\sigma_i^2 + \alpha^2} \approx \frac{\sigma_i^2}{\alpha^2} \ll 1$
- A “good” choice of α may provide enough numerical stability to expect a good approximate solution

³ The singular expression and the use of Tikhonov filters

Tikhonov regularization and Damped SVD

- A little algebra shows that the minimum solution of (5) is given by the nonsingular system

$$(A^T A + \alpha^2 I_n) x = A^T y$$

and from (4) we can show that

$$x^* = \sum_{i=1}^n f_i \left(\frac{u_i^T y}{\sigma_i} \right) v_i$$

where $f_i = \frac{\sigma_i^2}{\sigma_i^2 + \alpha^2}$ are known as *filter factors*³

- The impact of an small α in the filter factors is:
 - None for large σ_i ($\alpha \ll \sigma_i$), i.e. $\frac{\sigma_i^2}{\sigma_i^2 + \alpha^2} \approx 1$
 - Reduce the magnification of $\frac{1}{\sigma_i}$ since $\frac{\sigma_i^2}{\sigma_i^2 + \alpha^2} \approx \frac{\sigma_i^2}{\alpha^2} \ll 1$
- A “good” choice of α may provide enough numerical stability to expect a good approximate solution

³ The singular expression can be written as $\frac{\sigma_i^2}{\sigma_i^2 + \alpha^2} = \frac{\sigma_i^2}{\sigma_i^2 + \alpha^2}$

Tikhonov regularization order one

- Damping the large components in magnitude may not inhibit undesirable behavior of the singular values.
- Strong regularization is needed, penalizing rapid changes of x_i (4)

$$\min \frac{1}{2} \|Ax - y\|_2^2 + \frac{1}{2} \alpha^2 \sum_{i=2}^{n-1} (x_i - x_{i-1})^2$$

- Again this expression is minimized by the solution of

$$(A^T A + \alpha^2 B_1^T B_1) x = A^T y$$

where

$$B_1 = \begin{bmatrix} 1 & -1 & 0 & 0 & 0 \\ 0 & 1 & -1 & 0 & 0 \\ \vdots & \ddots & 1 & \ddots & \ddots \\ \vdots & \vdots & \vdots & \ddots & -1 \\ 0 & 0 & 0 & 0 & 1 \end{bmatrix}_{(n-1) \times n}$$

Tikhonov regularization order one

- Damping the large components in magnitude may not inhibit undesirable behavior of the singular values.
- Strong regularization is needed, penalizing rapid changes of x_i (4)

$$\min \frac{1}{2} \|Ax - y\|_2^2 + \frac{1}{2} \alpha^2 \sum_{i=2}^{n-1} (x_i - x_{i-1})^2$$

- Again this expression is minimized by the solution of

$$(A^T A + \alpha^2 B_1^T B_1) x = A^T y$$

where

$$B_1 = \begin{bmatrix} 1 & -1 & 0 & 0 & 0 \\ 0 & 1 & -1 & 0 & 0 \\ \vdots & \ddots & 1 & \ddots & \ddots \\ \vdots & \vdots & \vdots & \ddots & -1 \\ 0 & 0 & 0 & 0 & 1 \end{bmatrix}_{(n-1) \times n}$$

Tikhonov regularization order one

- Damping the large components in magnitude may not inhibit undesirable behavior of the singular values.
- Strong regularization is needed, penalizing rapid changes of x_i (4)

$$\min \frac{1}{2} \|Ax - y\|_2^2 + \frac{1}{2} \alpha^2 \sum_{i=2}^{n-1} (x_i - x_{i-1})^2$$

- Again this expression is minimized by the solution of

$$(A^T A + \alpha^2 B_1^T B_1) x = A^T y$$

where

$$B_1 = \begin{bmatrix} 1 & -1 & 0 & 0 & 0 \\ 0 & 1 & -1 & 0 & 0 \\ \vdots & \ddots & 1 & \ddots & \ddots \\ \vdots & \vdots & \vdots & \ddots & -1 \\ 0 & 0 & 0 & 0 & 1 \end{bmatrix}_{(n-1) \times n}$$

Tikhonov regularization order two

- An even stronger regularization is

$$\min \frac{1}{2} \|Ax - y\|_2^2 + \frac{1}{2} \alpha^2 \sum_{i=2}^{n-1} (x_{i+1} - 2x_i + x_{i-1})^2$$

- Again this expression is minimized by the solution of

$$(A^T A + \alpha^2 B_2^T B_2) x = A^T y$$

where

$$B_2 = \begin{bmatrix} -2 & 1 & 0 & 0 & \dots \\ 1 & -2 & 1 & 0 & \dots \\ \vdots & 1 & -2 & 1 & \vdots \\ \vdots & \vdots & \ddots & \ddots & \ddots \\ 0 & 0 & 0 & 1 & -2 \end{bmatrix}_{(n-2) \times n}$$

Tikhonov regularization order two

- An even stronger regularization is

$$\min \frac{1}{2} \|Ax - y\|_2^2 + \frac{1}{2} \alpha^2 \sum_{i=2}^{n-1} (x_{i+1} - 2x_i + x_{i-1})^2$$

- Again this expression is minimized by the solution of

$$(A^T A + \alpha^2 B_2^T B_2) x = A^T y$$

where

$$B_2 = \begin{bmatrix} -2 & 1 & 0 & 0 & \cdots \\ 1 & -2 & 1 & 0 & \cdots \\ \vdots & 1 & -2 & 1 & \vdots \\ \vdots & \vdots & \ddots & \ddots & \ddots \\ 0 & 0 & 0 & 1 & -2 \end{bmatrix}_{(n-2) \times n}$$

References



Numerical Optimization. J. Nocedal, S. Wright. Second Edition. Springer. 2006



Matrix Computations. G. Golub, Van Loan. Third Edition. Johns Hopkins University Press. 1996